Operating Systems
Paging

Lecture 10
Michael O'Boyle

Overview

- Paging
- Page Tables
- TLB
- Shared Pages
- Hierarchical Pages
- Hashed Pages
- Inverted Pages
- Uses

Address Translation Scheme

■ Address generated by CPU is divided into:

- **Page number** (*p*) – used as an index into a **page table** which contains base address of each page in physical memory
- **Page offset** (*d*) – combined with base address to define the physical memory address that is sent to the memory unit

page number	page offset
p	d
m -n	n

- For given logical address space 2^m and page size 2^n

Paging Hardware
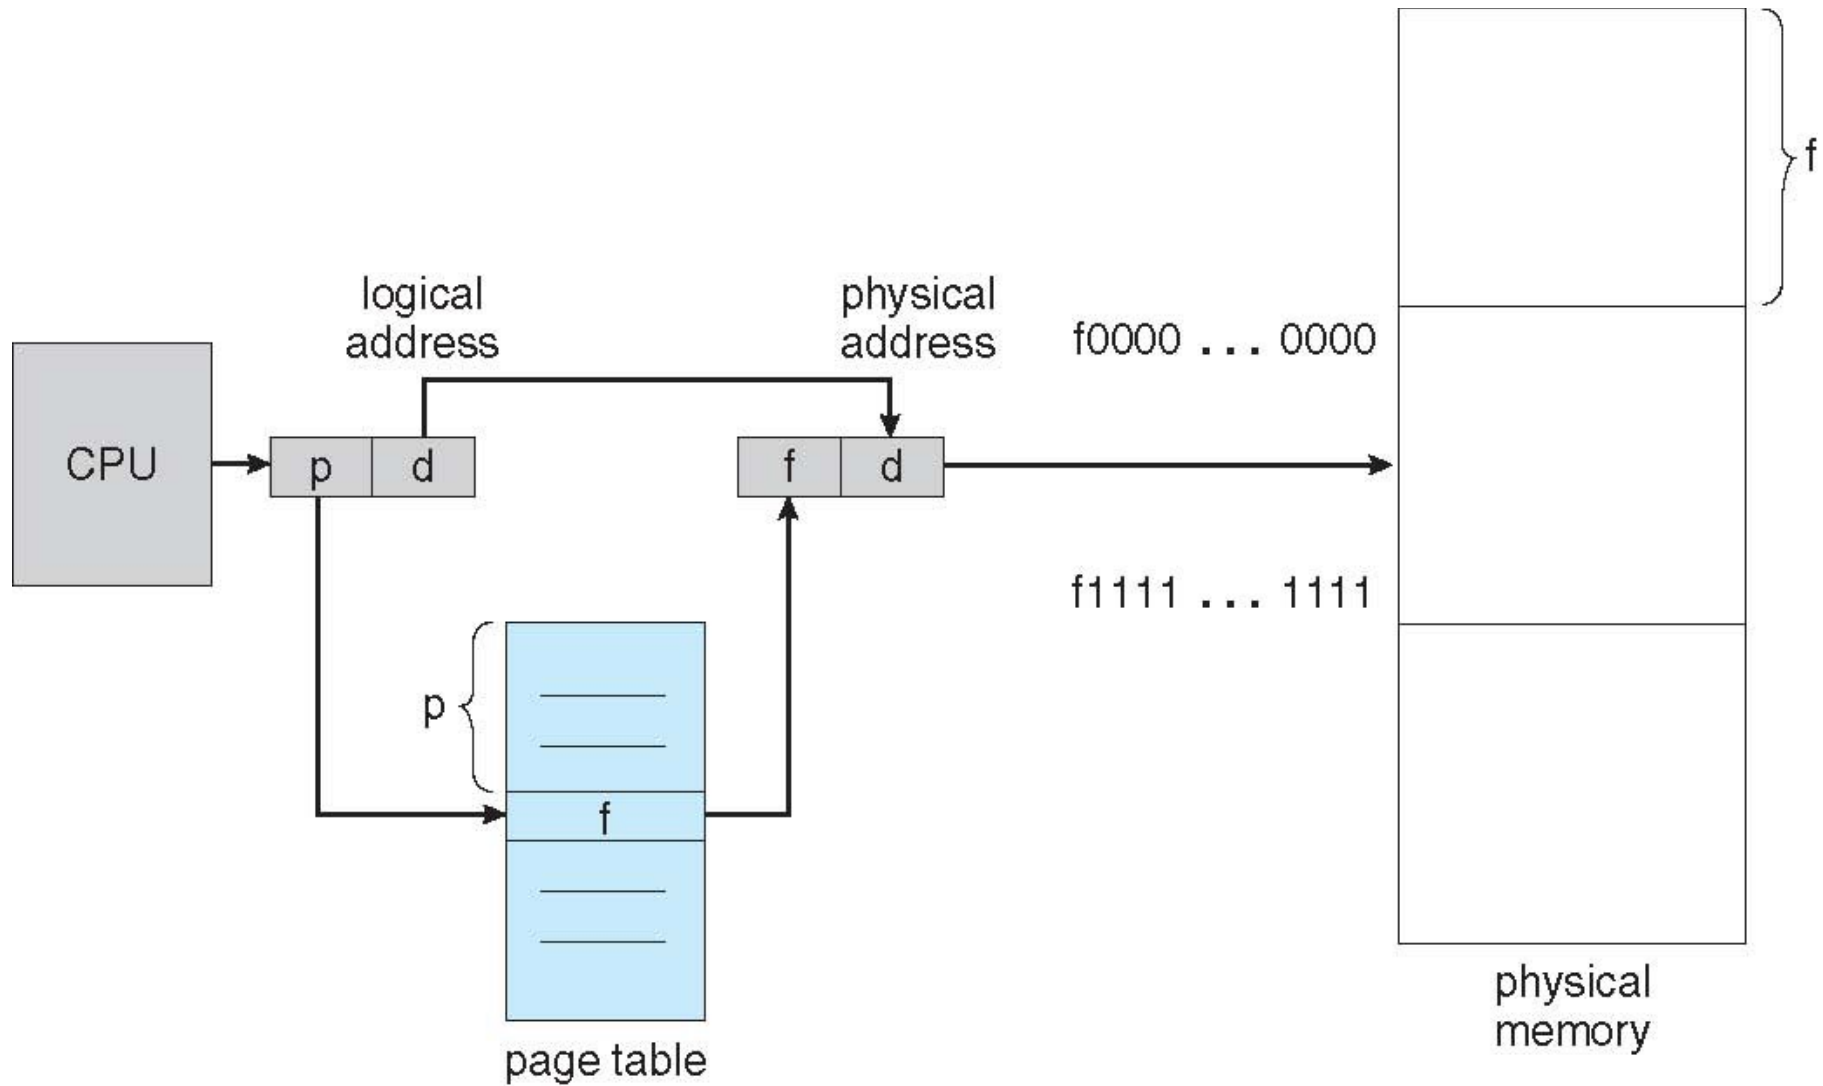
Paging Model of Logical and Physical Memory

page 0
page 1
page 2
page 3

logical
memory

0	1
1	4
2	3
3	7

page table

frame
number

0	
1	page 0
2	
3	page 2
4	page 1
5	
6	
7	page 3

physical
memory

Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

|	
0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

*n*=2 and *m*=4   32-byte memory and 4-byte pages

Paging (Cont.)

- Calculating internal fragmentation
  - Page size = 2,048 bytes
  - Process size = 72,766 bytes
  - 35 pages + 1,086 bytes
  - Internal fragmentation of 2,048 - 1,086 = 962 bytes
  - Worst case fragmentation = 1 frame – 1 byte
  - On average fragmentation = 1 / 2 frame size
  - So small frame sizes desirable?
  - But each page table entry takes memory to track
  - Page sizes growing over time
    - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

Free Frames

Before allocation (a)

After allocation (b)

Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register** (**PTBR**) points to the page table
- **Page-table length register** (**PTLR**) indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
  - One for the page table and one for the data / instruction
- The two memory access problem can be solved
  - by the use of a special fast-lookup hardware cache
  - called **associative memory** or **translation look-aside buffers** (**TLBs**)

Implementation of Page Table

- Some TLBs store **address-space identifiers** (**ASIDs**) in each TLB entry –
  - uniquely identifies each process
  - provide address-space protection for that process
  - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
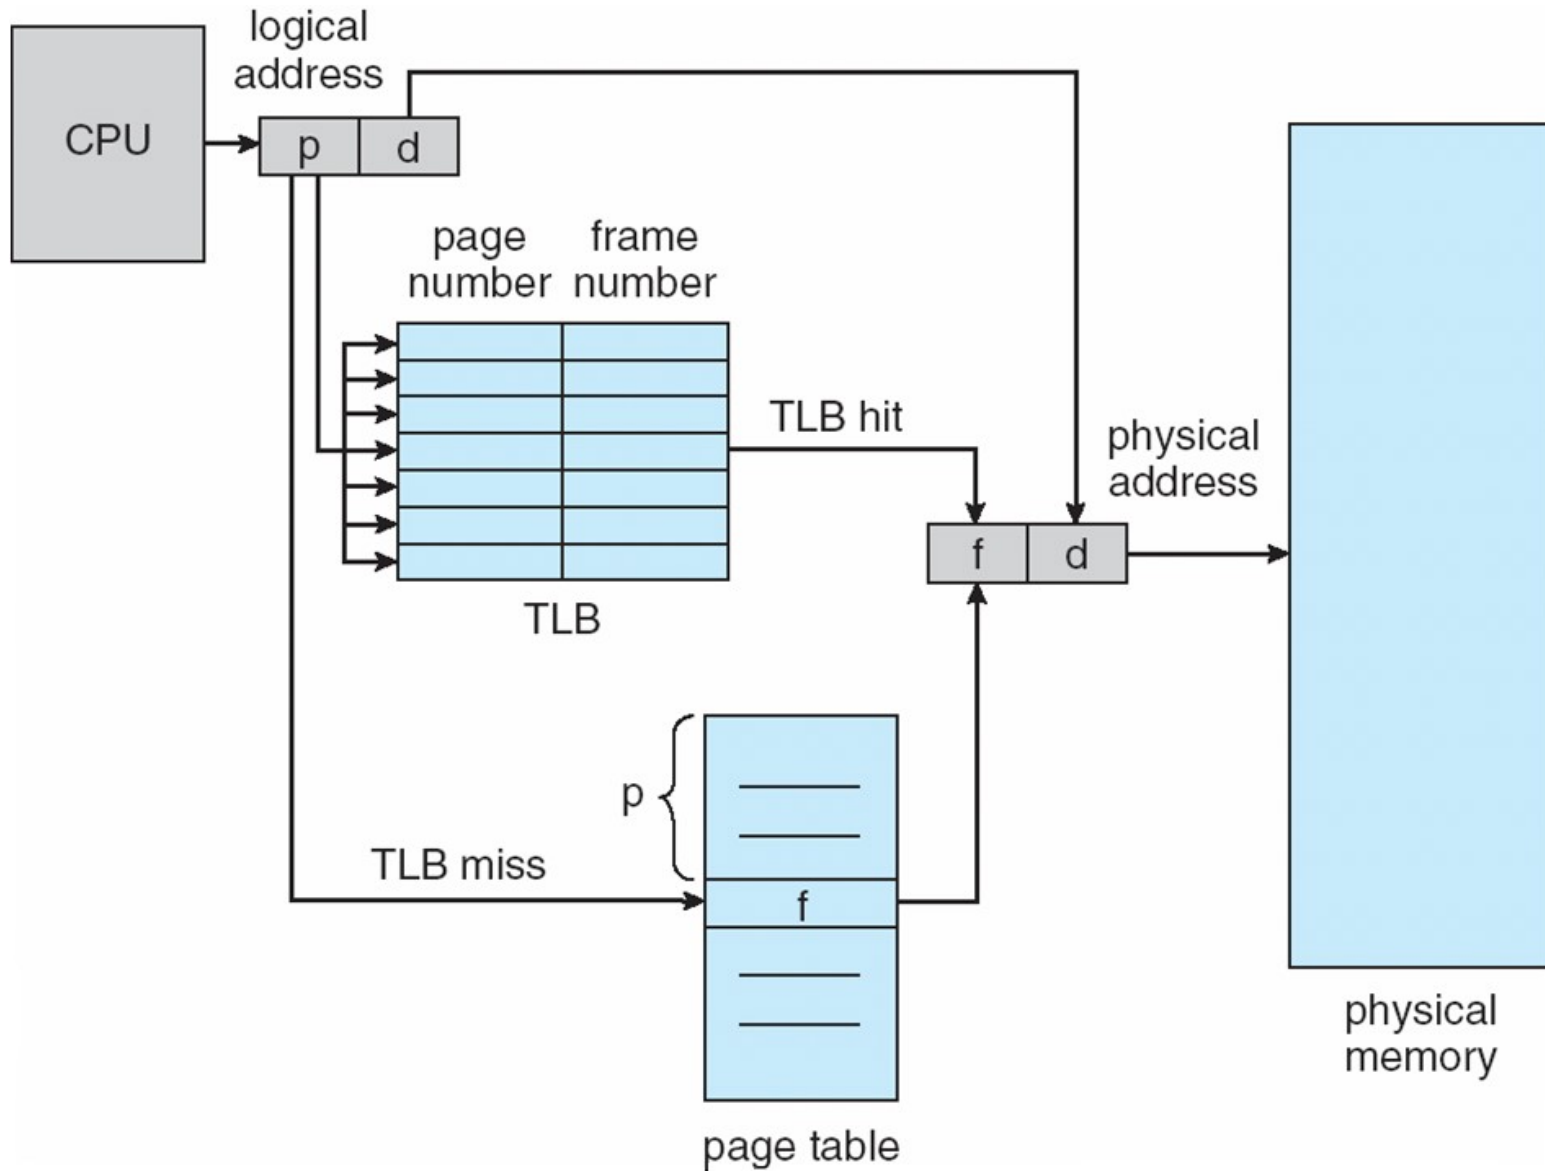 - Some entries can be **wired down** for permanent fast access

Associative Memory

- Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory

Paging Hardware With TLB

Effective Access Time

- ## Associative Lookup
  - Extremely fast

- ## Hit ratio = α
  - Hit ratio – percentage of times that a page number is found in the associative memory ;
  - Consider α = 80%, 100ns for memory access

- ## Consider α = 80%, 100ns for memory access
  - EAT = 0.80 x 100 + 0.20 x 200 = 120ns

- ## Consider hit ratio α = 99, 100ns for memory access
  - EAT = 0.99 x 100 + 0.01 x 200 = 101ns

Memory Protection

- Memory protection implemented
  - by associating protection bit with each frame
  - to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - "valid" indicates that the associated page
    - is in the process' logical address space, and is thus a legal page
  - "invalid" indicates that the page I
    - is not in the process' logical address space
  - Or use **page-table length register** (**PTLR**)
  - Page Table Entries (PTEs) can contai more information
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table

frame number   valid–invalid bit

0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

00000

page 0
page 1
page 2
page 3
page 4

10,468
12,287
page 5

0
1
2 page 0
3 page 1
4 page 2
5
6
7 page 3
8 page 4
9 page 5
   ...
   page n

Shared Pages

- **Shared code**
  - One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
  - Similar to multiple threads sharing the same process space
  - Also useful for interprocess communication if sharing of read-write pages is allowed

- **Private code and data**
  - Each process keeps a separate copy of the code and data
  - The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example
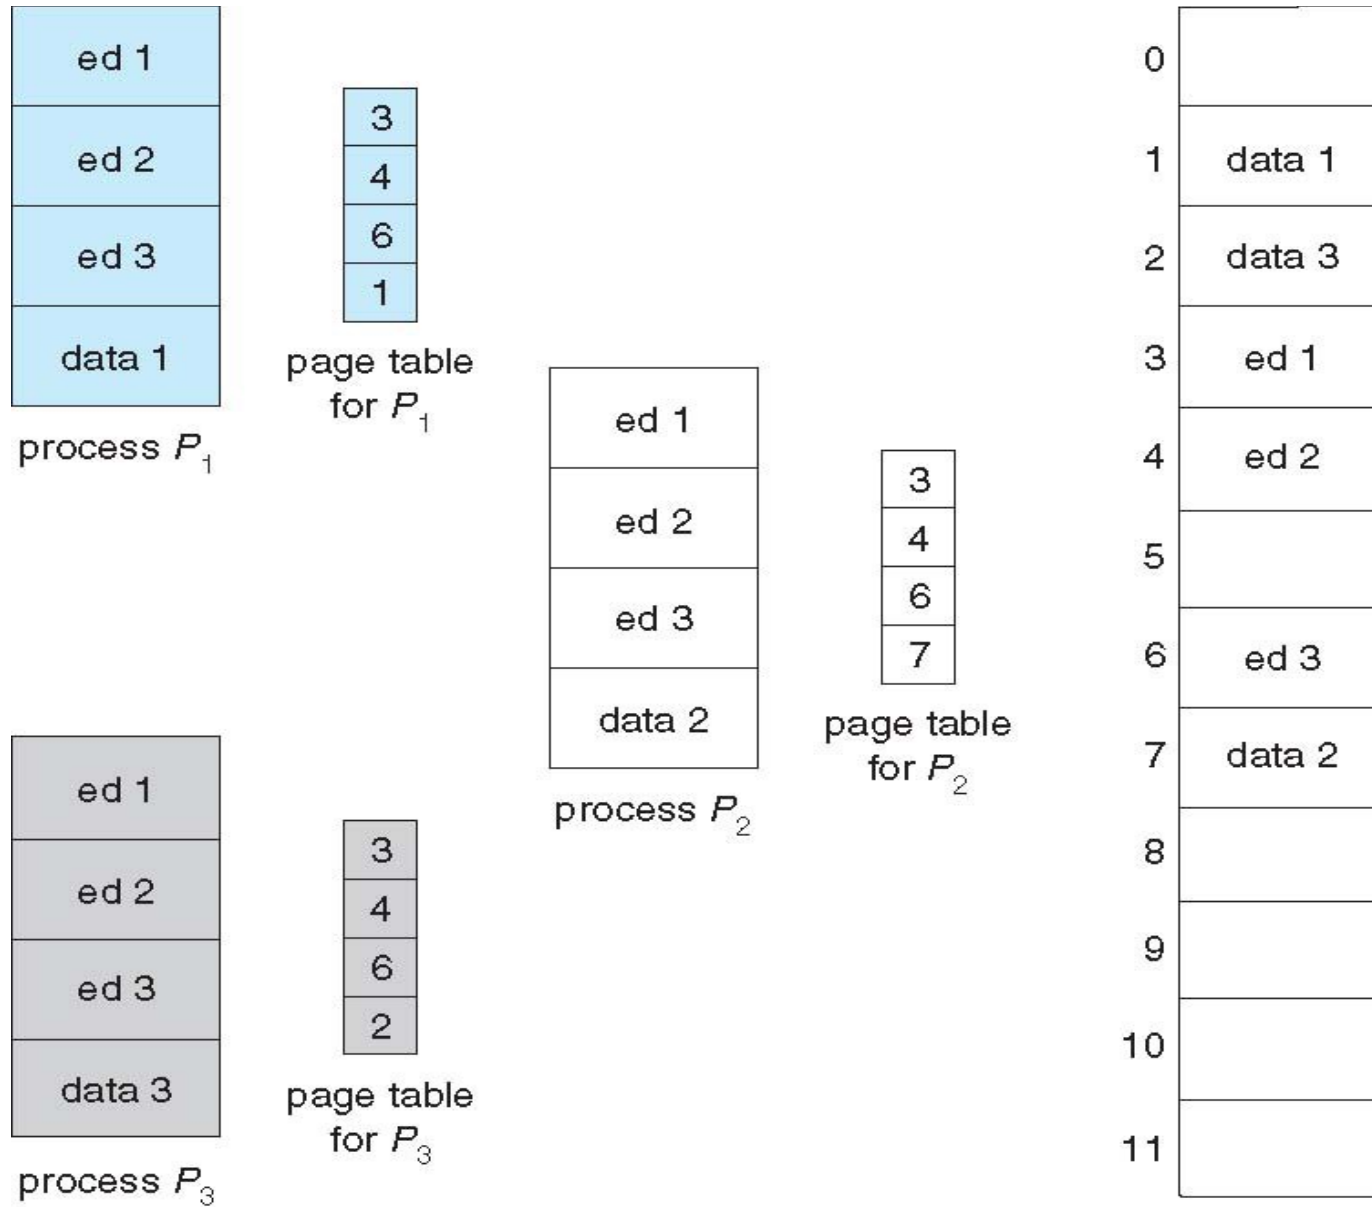
Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
    - Consider a 32-bit logical address space as on modern computers
    - Page size of 4 KB (2^{12})
    - Page table would have 1 million entries ($2^{32} / 2^{12}$)
    - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
        - That amount of memory used to cost a lot
        - Don't want to allocate that contiguously in main memory
- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

Hierarchical Page Tables

- Break up the logical address space into multiple page tables

- A simple technique is a two-level page table

- We then page the page table
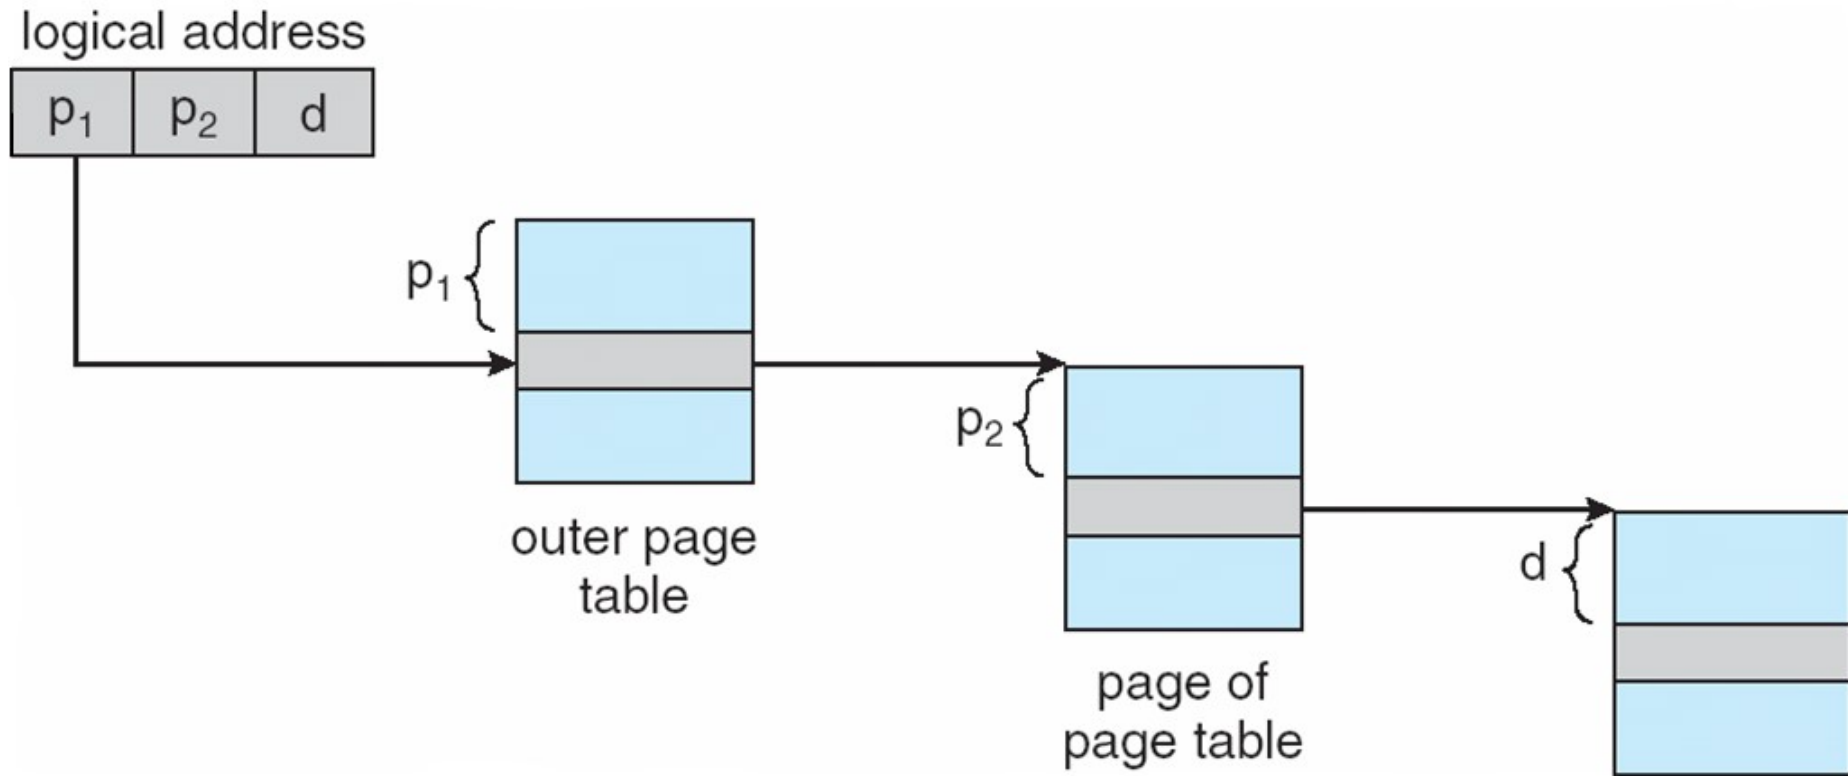
Two-Level Page-Table Scheme
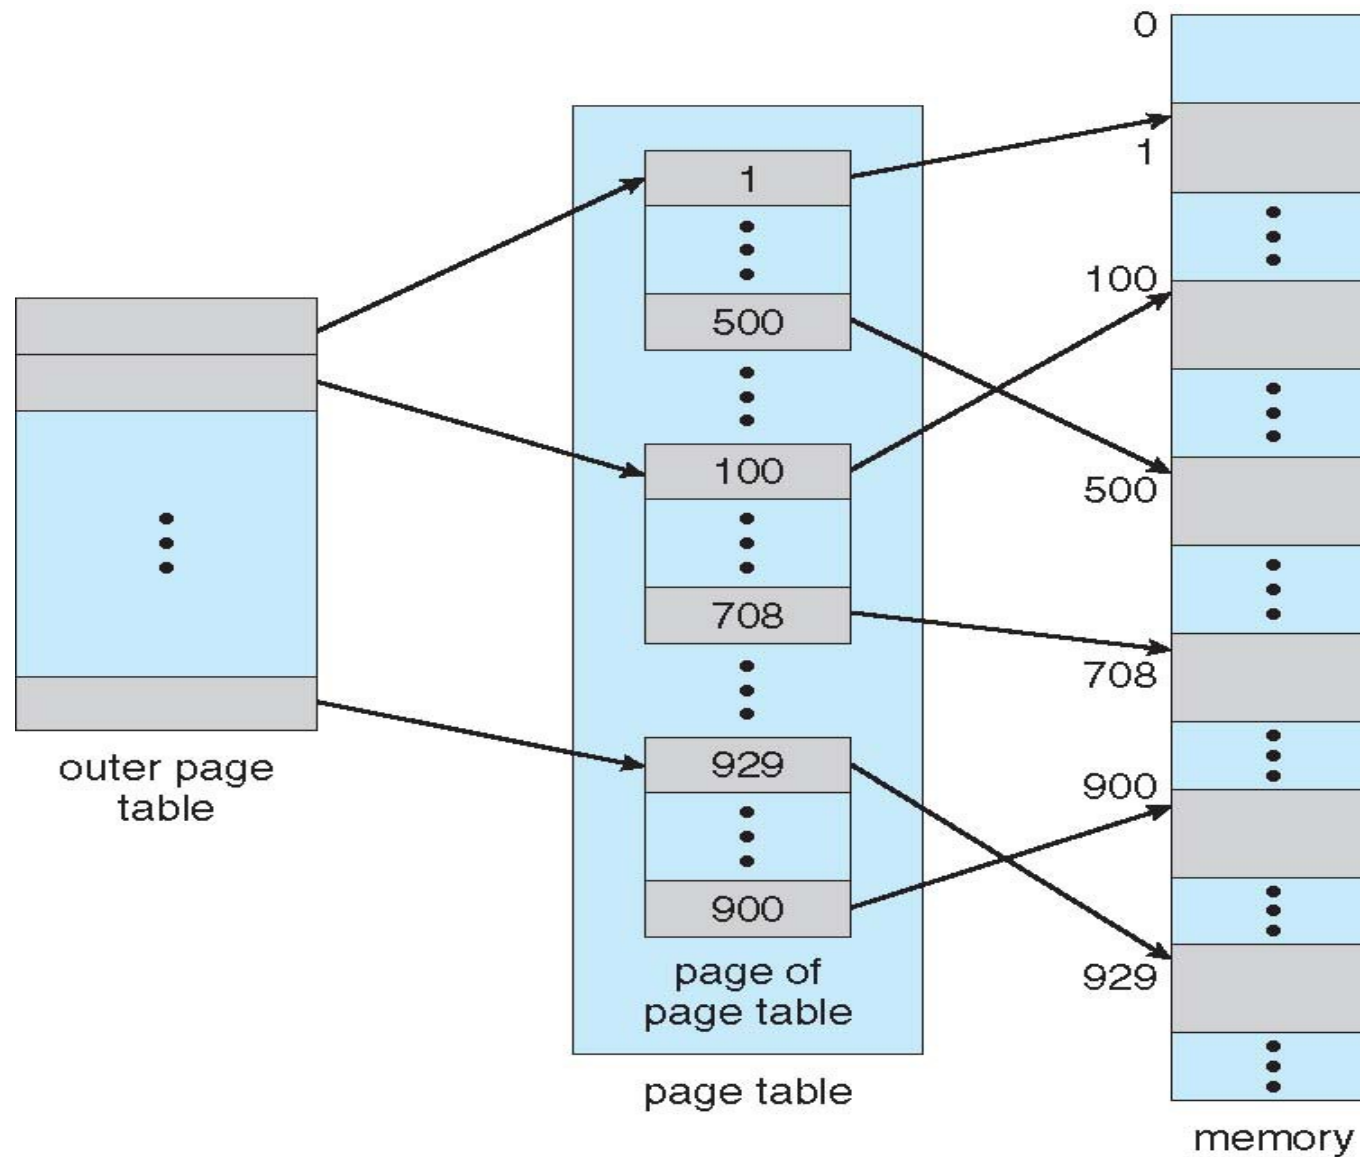
Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits

- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset

- Thus, a logical address is as follows:

| page number		page offset
p_1	p_2	d
12	10	10

- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

Address-Translation Scheme

64-bit Logical Address Space

■ Even two-level paging scheme not sufficient

■ If page size is 4 KB (2^{12})

- Then page table has 2^{52} entries
- If two level scheme, inner page tables could be 2^{10} 4-byte entries
- Address would look like

outer page	inner page	page offset
p_1	p_2	d
42	10	12

- Outer page table has 2^{42} entries or 2^{44} bytes
- One solution is to add a 2nd outer page table
- But in the following example the 2nd outer page table is still 2^{34} bytes in size
  ‣ And possibly 4 memory access to get to one physical memory location

Three-level Paging Scheme

outer page	inner page	offset
p_1	p_2	d
42	10	12

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

Hashed Page Tables

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains
  - (1) the virtual page number
  - (2) the value of the mapped page frame
  - (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
  - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)

Hashed Page Table

Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages,
    - track all physical pages
- One entry for each real page of memory
- Entry consists of
    - the virtual address of the page stored in that real memory location,
    - information about the process that owns that page
- Decreases memory needed to store each page table
    - but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one/few page-table entries
    - TLB can accelerate access
- But how to implement shared memory?
    - One mapping of a virtual address to the shared physical address

Inverted Page Table Architecture
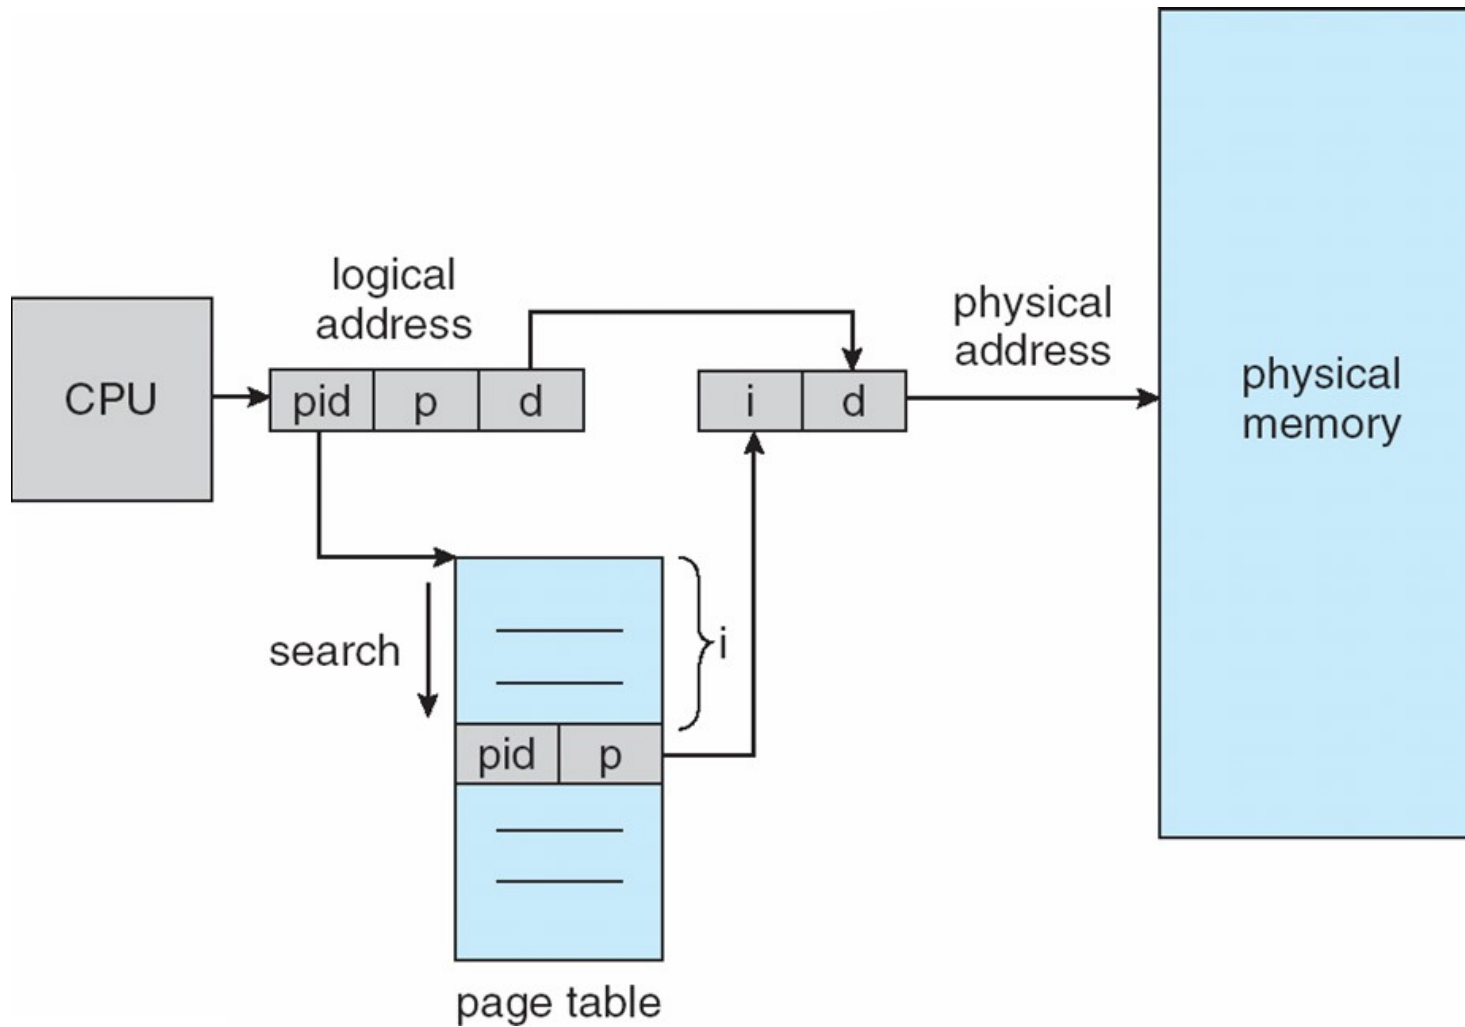
Functionality enhanced by page tables

- Code (instructions) is read-only
  - A bad pointer can't change the program code
- Dereferencing a null pointer is an error caught by hardware
  - Don't use the first page of the virtual address space – mark it as invalid – so references to address 0 cause an interrupt
- Inter-process memory protection
  - My address XYZ is different that your address XYZ
- Shared libraries
  - All running C programs use libc
  - Have only one (partial) copy in physical memory, not one per process
  - All page table entries mapping libc point to the same set of physical frames
    - DLL's in Windows

More functionality

- Generalizing the use of "shared memory"
  - Regions of two separate processes's address spaces map to the same physical frames
  - Faster inter-process communication
    - Just read/write from/to shared memory
    - Don't have to make a syscall
  - Will have separate Page Table Entries (PTEs) per process, so can give different processes different access rights
    - E.g., one reader, one writer
- Copy-on-write (CoW), e.g., on fork()
  - Instead of copying all pages, create shared mappings of parent pages in child address space
    - Make shared mappings read-only for both processes
    - When either process writes, fault occurs, OS "splits" the page

Less familiar uses

- Memory-mapped files
    - instead of using open, read, write, close
        - "map" a file into a region of the virtual address space
            - e.g., into region with base 'X'
        - accessing virtual address 'X+N' refers to offset 'N' in file
        - initially, all pages in mapped region marked as invalid
    - OS reads a page from file whenever invalid page accessed
    - OS writes a page to file when evicted from physical memory
        - only necessary if page is dirty

More unusual use

- Use "soft faults"
  - faults on pages that are actually in memory,
  - but whose PTE entries have artificially been marked as invalid
- That idea can be used whenever it would be useful to trap on a reference to some data item
- Example: debugger watchpoints
- Limited by the fact that the granularity of detection is the page

Summary

- Paging
- Page Tables
- TLB
- Shared Pages
- Hierarchical Pages
- Hashed Pages
- Inverted Pages
- Uses
- Next time: Virtual Memory